



# **Deployment of Autopilot Car Sharing Service and requirements within IBM Cloud Cluster**

## **Project Report**

from the Course of Studies Applied Computer Science  
at the Cooperative State University Baden-Württemberg Mannheim

by  
**Pascal Schroeder**

11/03/2019

**Time of Project**  
**Student ID, Course**  
**Division, Business Unit**  
**Company**  
**Supervisor in the Company**

19.11.2018-08.02.2019  
5501463, TINF16AI-BC  
IBM Research  
IBM Deutschland GmbH, Dublin  
Martin Mevissen

**Signature Supervisor**

---

# Confidentiality Statement

The Project Report on hand

*Deployment of Autopilot Car Sharing Service and requirements within IBM Cloud Cluster*

contains internal resp. confidential data of IBM Deutschland GmbH. It is intended solely for inspection by the assigned examiner, the head of the Applied Computer Science department and, if necessary, the Audit Committee at the Cooperative State University Baden-Württemberg Mannheim. It is strictly forbidden

- to distribute the content of this paper (including data, figures, tables, charts etc.) as a whole or in extracts,
- to make copies or transcripts of this paper or of parts of it,
- to display this paper or make it available in digital, electronic or virtual form.

Exceptional cases may be considered through permission granted in written form by the author and IBM Deutschland GmbH.

Mannheim, 11/03/2019

---

Pascal Schroeder

# Contents

<b>Declaration of Sincerity</b>	<b>I</b>
<b>Abstract</b>	<b>II</b>
<b>1 Introduction</b>	<b>1</b>
1.1 Introduction to ‘Autopilot’ project . . . . .	1
1.2 Advantage of cluster systems . . . . .	2
1.3 Project objective . . . . .	3
<b>2 Theory</b>	<b>4</b>
2.1 Kubernetes cluster solution . . . . .	4
2.2 Apache Kafka . . . . .	9
2.3 Open Source Routing Machine . . . . .	9
2.4 IoT Platform . . . . .	9
2.5 DRL Car Sharing Service . . . . .	9
<b>3 Method</b>	<b>10</b>
3.1 Catalogue of Criteria . . . . .	10
3.2 Locally setup of car sharing app . . . . .	10
3.3 Outsourcing kafka client to IBM Cloud message hub . . . . .	10
3.4 Dockerizing car sharing app and requirements . . . . .	10
3.5 Deploying and exposing docker container on cluster . . . . .	10
<b>4 Result</b>	<b>11</b>
4.1 Outsourcing of Apache Kafka Client . . . . .	11
4.2 Communication between car sharing app and services . . . . .	11
<b>5 Discussion</b>	<b>12</b>
<b>Acronyms</b>	<b>III</b>
<b>Bibliography</b>	<b>IV</b>
<b>Appendices</b>	<b>VII</b>

# Declaration of Sincerity

Hereby I solemnly declare that my project report, titled *Deployment of Autopilot Car Sharing Service and requirements within IBM Cloud Cluster* is independently authored and no sources other than those specified have been used.

I also declare that the submitted electronic version matches the printed version.

Mannheim, 11/03/2019

Place, Date

---

Signature Pascal Schroeder



# Abstract

# 1 Introduction

## 1.1 Introduction to 'Autopilot' project

The AUTOPILOT project is a Large-Scale Pilot (LSP) by the European Commission. Those LSPs include five pilots as well as two coordination and support actions. Thereby the AUTOPILOT project takes care of using Internet of Things (IoT) for connecting autonomous driving to a larger environment like traffic monitoring systems, car parks, traffic light radars or just other vehicles.

In this project 45 different partners from 15 different countries across Europe are working together in four different sectors - development of autonomous driving vehicles, development of IoT and networks, collection of data to evaluate the systems and their potential impacts and organisations that will use the results for developing innovative services.

This project not only connects different sensors, cameras and vehicles to the IoT Platform, but also includes autonomous driving modes like automated valet parking, car sharing or a highway pilot, as well as driving services like route optimisation, chauffeur services or electronic driving licenses and many more.

In the IBM Research Lab the Optimization & Control team is working on implementing different IoT platforms and networks for creating a large environment with many vehicles, road sensors and infrastructure facilities. This solves the limitations usual on-board sensors cause, because it expands the point of view from the vehicles view only to a far larger environment. This approach enables possibilities like anticipating situations upfront, minimize risks like traffic jam or accidents and all in all improving safety and comfort while reducing costs through a better organized and connected traffic.

One part of this is a car sharing service for autonomous driving vehicles including a scheduler, which matches the customers to its target and calculates the best fitting vehicle and optimizes the route. Thereby it also includes traffic information and recalculate the route dependent on traffic jams, accidents or other environmental changes.

For proving the functionality of this scheduler it is necessary to test it in different environments - first in a real-life environment with working, autonomous driving cars and second in a large scalable simulation for proving its functionality in a large environment with many cars and customers.

This also needs to be based on a reliable and scalable infrastructure with enough resources to execute all the necessary calculations. One modern approach for this are cluster systems, of which advantages will be described in chapter 1.2.

## 1.2 Advantage of cluster systems

Cluster systems are two or more computers connected to each other, so that they can share their resources and can be viewed as one system. Therefore they can be connected physically as well as virtually. Using such a cluster system is typically much more cost-efficient compared to using a single computer, which provides about the same power. But the advantage of such systems doesn't end with a better cost-efficiency.<sup>cmp.[1]</sup>

First clusters are not only used for one specific purpose, but there are different kinds of cluster configurations, which are build for different objectives.

On the one hand, there are "Load balancing" cluster configurations. Load balancing clusters are clusters with the objective to provide better computing performance, like minimizing response time and maximizing throughput. Therefore it splits the computation in different parts and runs them parallel on different nodes. Through that the capabilities of the systems can be combined.<sup>cmp.[2]</sup>

Furthermore there are "High availability" clusters, which are designed to prevent a total break down of the system. For that there are several redundant nodes connected to the system. That means, that when one component fails, a redundant node of this component is used to provide the service. Thereby availability is ensured even if one component fails. The more redundant nodes are used the better will be the availability of a system. To provide an example: If a system has an availability of 90% it has a downtime of 10%, which means that the system is not available on 36.53 days per year. If there is a second, redundant and independent node connected to the cluster, which acts as a stand-in for the first system in case of a failure, the downtime of those two systems has to be combined. That means that this cluster would only fail if there is an overlap between the downtimes of those two systems. Resulting of a downtime of 10% of each system there is a combined downtime of 1% or 3.65 days per year. If there is a third system connected to the cluster, the downtime will be reduced to 0.1% or 8.77 hours per year and so on. The job of "High availability" clusters is in that to reduce the downtime of the system and to provide a working system almost any time.<sup>cmp.[3]</sup>

Figure 1.1 shows an example architecture of such a "High availability" cluster. There are two servers, whereby one server is a replication of the other one. The request router handles the incoming requests and leads them to the active server. If there is a system




Figure 1.1.1 “High availability” cluster architecture

failure of the active server, the manager will recognize that, and sends a message to the request router to change the request address to the other server. From now on this server will deal with all the requests. Also there is a shared storage for both of the servers in one database. No matter which one is active they can both access this database.<sup>cmp.[4]</sup>

But those advantages don't have to stand alone, but they can also be combined, so that you can provide high availability as well as high performance. Because there is an continuously increasing amount of calculations to be made cluster systems are becoming more and more important in modern IT for providing a solution for both problems.

## 1.3 Project objective



## 2 Theory

### 2.1 Kubernetes cluster solution

One very common system for managing cluster systems as described in chapter 1.1 is Kubernetes - or short **K8s**. Originally developed by Google and now maintained by the Cloud Native Computing Foundation, Kubernetes is an “open-source platform for managing containerized workloads and services”.<sup>cmp.[12]</sup> For understanding what that means the concept of containers needs to be described first.

Containers are isolated, stand-alone packages of software, similar to processes. In those packages everything is included, which this piece of software needs, like runtime, libraries, settings and other system tools. These containers have a completely different environment within themselves than outside. This environment includes for example network routes, dns settings and control group limits. This enables the possibility to share common resources and still be isolated from any other process as well as the host system. Thereby containers are always working the same, no matter on what system they run or in which environment.<sup>cmp.[13], [14]</sup>

Kubernetes is for an automating deployment, scaling and management of these containers within a cluster of nodes. Thereby a cluster consists of at least one master node and any number of worker nodes. Figure 2.1.1 shows the different services owned by master and worker nodes.

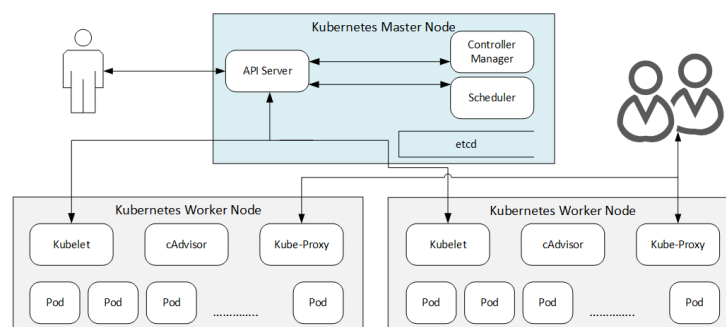


Figure 2.1.1 Kubernetes allocation of services  
Based on retrieved information from [13]

First there are several pods on each worker node. Pods are the smallest unit in Kubernetes. They contain one or more containers, which are deployed together on the same host. Thereby they can work together to perform a set of tasks.<sup>cmp.[15]</sup>

On the master node there are an **API** (Application Programming Interface) Server, a Controller Manager, a Scheduler and a key-value store called `etcd`.<sup>cmp.[13], [16]</sup>

The API Server is for clients to run their requests against. That means the API Server is responsible for the communication between Master and Worker nodes and for updating corresponding objects in the `etcd`. Also the authentication and authorization is task of the API Server. The protocol for the communication is written in **REST** (Representational State Transfer). For reacting on changes of clients there is also a watch mechanism implemented, which triggers an action after some specific changes, like the scheduler creating a new pod. This workflow is showed in figure 2.1.2.<sup>cmp.[13], [16]</sup>

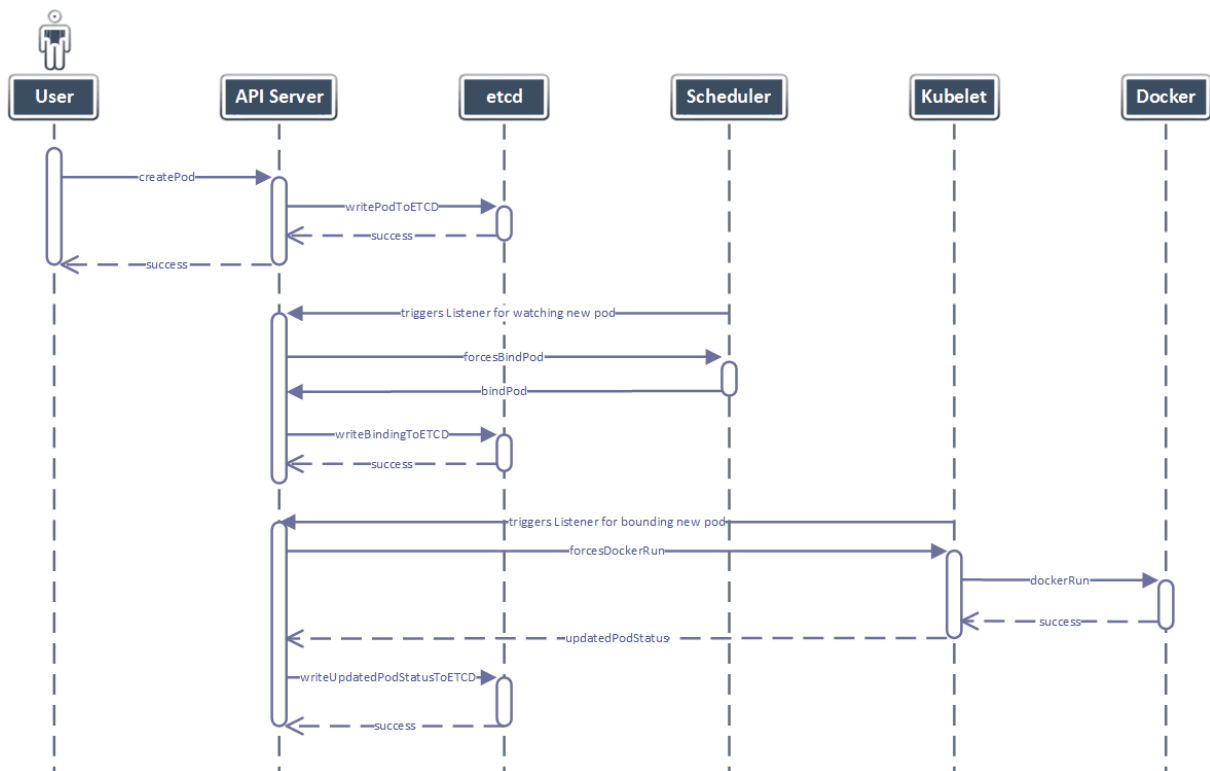


Figure 2.1.2 Kubernetes API Server watcher sequence diagram

Based on retrieved information from [16]

There you can see the user creating a pod through a belonging request to the API Server. The API Server writes this change to the `etcd`. Afterwards the API Server recognizes a new pod in the `etcd` and invokes the Scheduler to create this new pod. What is the exact task of the scheduler is described later. After successfully creating and binding the new pod, the API Server writes this change to the `etcd`. Because of this change within the `etcd` the API Server invokes the kubelet, which is also described later in this chapter, of the corresponding node. This kubelet starts docker to create the containers of the pod. Kubelet responses with the new status of the pod, which is then written to the `etcd` by the API Server. After that the creation of the new pod is successfully finished.<sup>cmp.[13], [16]</sup>

The Controller Manager is a daemon, which embeds all of the Kubernetes controller. Examples for them are the Replication Controller or the Endpoint Controller. Those controllers are watching the state of the cluster through the API Server. Whenever a specific action happens, it performs the necessary actions to hold the current state or to move the cluster towards the desired state. For providing an example: If the Replication Controller recognizes, that one replication of a pod has been destroyed for some reason, it will take care of triggering the creation of a new replication.<sup>cmp.[13], [16]</sup>

The scheduler manages the binding of pods to nodes. Therefore it watches for new deployments as well as for old ones to create new pods if a new deployment is created or recreating a pod whenever a pod gets destroyed. The scheduler organizes the allocation of the pods within the cluster on the basis of available resources of the pods. That means, that it always create pods, where the most resources are available, or reorganize the allocation if there is a change in the resource allocation of the cluster.<sup>cmp.[13], [16]</sup>

Figure 2.1.3 shows the way the Scheduler works, when new nodes are connected. As long as there are only two nodes and four pods need to be deployed, it allocates those four pods to the two nodes. As soon as there are more nodes connected, the scheduler recognizes free resources and reallocate the pods to these new nodes.




Figure 2.1.3 Kubernetes API Server watcher sequence diagram

Based on retrieved information from [17]

The etcd is a key-value store, which stores the configuration data and the condition of the Kubernetes cluster. The etcd also contains a watch feature, which listens to changes to keys and triggers the API server to perform all necessary actions to move the current state of the cluster towards the desired state.<sup>cmp.[13], [16]</sup>

The worker node consists of a Kubelet, a cAdvisor, a Kube-Proxy and - as mentioned before - several Pods.

The kubelet needs to be used if a new pod should be deployed. Then it gets the action to create all needed containers. For that it uses Docker to create them. Afterwards it combines some containers into one pod. Containers in one pod are always started and stopped together. This pod will then be deployed on the node, on which the kubelet is located.<sup>cmp.[13], [16]</sup>

The cAdvisor measures the usage of CPU-resources as well as demanded memory on the node, on which it is located, and notifies the master about it. Based on those measurements the scheduler allocates the pods within the cluster to ensure the best possible allocation of resources.<sup>cmp.[13], [16]</sup>

The kube-proxy is a daemon, that runs as a simple network proxy to provide the possibility of communicating to that node within the cluster. Additionally it runs a load balancer for the services on that node.<sup>cmp.[13], [16]</sup>

Through this architecture Kubernetes enables different possibilities to deploy pods within the cluster. The simplest one is to deploy a specific pod directly through the

```
1 kubectl create -f 'image\_path'
```

Listing 2.1: Create Kubernetes pod

command. This deploys the pod described in the file, but it doesn't ensure the failure safety. That means if the pod gets destroyed for some reason, no matter if it has been destroyed by accident or intended, it won't be recreated and deployed again automatically.

Another possibility is to create a deployment. Therefore the image has to be embedded within a replicaset and this replicaset within a deployment. If this deployment is created, it will automatically create every pod of the replicaset. If one pod gets terminated, no matter if manually or because of an error, it will be directly recreated and deployed.<sup>cmp.[13], [18]</sup>

This procedure also enables the possibility to execute dynamical rolling updates. If you execute the command

```
1 kubectl set image deployment/name-deployment name=name:1.1.1
```

Listing 2.2: Create Kubernetes deployment

Kubernetes will start a Rolling Update. This causes the creation of a new replicaset, which uses pods of the new version. While the new replicaset will be scaled up, the old one will be scaled down step by step. This enables the maintenance of a deployment even while an update is enrolled, so that the users can still use the services while those are updating.

That means, that there is never a need of a system to shut down because of an update, which needs to be enrolled, and the system guarantees its high availability.<sup>cmp.[13], [18]</sup>

A third possibility to deploy pods are services. Services are used to enable the usage of pods from outside the cluster. Therefore it gets the pod from the targetPort of the belonging node and creates a random port on this node. This port serves as endpoint for the LoadBalancer. Through this, everybody can now communicate with this pod.<sup>cmp.[13], [18]</sup>

This is how Kubernetes ensures High Availability as well as Load Balancing at the same time. Through the possibility of replicate every pod several times it is ensured that whenever one pod fails for some reason, another is already prepared to help out and take over the job. Even the master can and should be replicated to ensure a functional system, even if one master has been destroyed. For production environment it's recommended to have 3, 5 or 7 master nodes running at the same time. For facilitating the High Availability almost every unit of a Kubernetes cluster is stateless and could be run redundant in several instances. Only the etcd key value store is stateful. For solving that problem in case of failure of the node, which owns the etcd, there is a leader election between all master node replicas to determine the new leading etcd.<sup>cmp.[13], [18]</sup>

All in all Kubernetes combines all the benefits of cluster applications in one software. High availability as well as load balancing is guaranteed and it also ensures a high scalability, rolling updates and auto-scaling. Compared to other cluster systems, like Docker swarm, it has the biggest community and it can support clusters of up to 5000 nodes, which enables using Kubernetes even for very big clusters. For example the whole Google Cloud System runs on a Kubernetes cluster. Performing tests of Kubernetes shows, that in 99% the API responses in less than one second and also their pods start in less than five seconds in 99% when starting containers with pre pulled images. This all makes Kubernetes maybe one of the most flexible and fastest cluster systems. In exchange for that flexibility it is more difficult to set it up, because you have to do it all on your own. But only that way it can be ensured, that it runs fast and flexible in the same time, which is the reason for Kubernetes being so popular.<sup>cmp.[19], [20]</sup>



## **2.2 Apache Kafka**

## **2.3 Open Source Routing Machine**

## **2.4 IoT Platform**

## **2.5 DRL Car Sharing Service**

## **3 Method**

### **3.1 Catalogue of Criteria**

### **3.2 Locally setup of car sharing app**

### **3.3 Outsourcing kafka client to IBM Cloud message hub**

### **3.4 Dockerizing car sharing app and requirements**

### **3.5 Deploying and exposing docker container on cluster**

## **4 Result**

### **4.1 Outsourcing of Apache Kafka Client**

### **4.2 Communication between car sharing app and services**

Those results will be discussed in chapter 5.1



## 5 Discussion

# Acronyms

**K8s** Kubernetes

**API** Application Programming Interface

**REST** Representational State Transfer

# Bibliography

- [1] David A. Bader, Robert Pennington (2001): The International Journal of High Performance Computing Applications, Retrieved from [www.cc.gatech.edu/~bader/papers/ijhpca.pdf](http://www.cc.gatech.edu/~bader/papers/ijhpca.pdf)
- [2] Aho, Alfred V.; Blum, Edward K. (2011): Computer Science: The Hardware, Software and Heart of It,
- [3] Patent by Omar M. A. Gadir, Kartik Subbanna, Ananda R. Vayyala, Hariprasad Shanmugam, Amod P. Bodas, Tarun Kumar Tripathy, Ravi S. Indurkar, Kurma H. Rao for NetAppInc (07-23-2001): High-availability cluster virtual server system, US6944785B2, Retrieved from [patents.google.com/patent/US6944785B2/en](http://patents.google.com/patent/US6944785B2/en)
- [4] OV ( 09-21-2017) - Bitbucket: High availability for Bitbucket Server, Retrieved from [confluence.atlassian.com/bitbucketserver/high-availability-for-bitbucket-server-776640137.html](http://confluence.atlassian.com/bitbucketserver/high-availability-for-bitbucket-server-776640137.html)
- [5] Erik Thomsen (2002): OLAP Solutions: Building Multidimensional Information Systems, John Wiley & Sons inc., Page 5-8
- [6] Anja Bog (2014): Benchmarking Transaction and Analytical Processing Systems: The Creation of a Mixed Workload Benchmark and its Application, Springer, Page 1 Retrieved from [console.bluemix.net/docs/services/conversation/index.html#about](http://console.bluemix.net/docs/services/conversation/index.html#about)
- [7] Dr. Andreas Weininger (01-23-2018) HTAP – Hybrid Transactional and Analytics Processing, Retrieved from [www.informatik-aktuell.de/betrieb/datenbanken/htap-hybrid-transactional-and-analytics-processing.html](http://www.informatik-aktuell.de/betrieb/datenbanken/htap-hybrid-transactional-and-analytics-processing.html)
- [8] Cameron Coles (2018): 11 Advantages of Cloud Computing and How Your Business Can Benefit From Them, Retrieved from: [www.skyhighnetworks.com/cloud-security-blog/11-advantages-of-cloud-computing-and-how-your-business-can-benefit-from-them/](http://www.skyhighnetworks.com/cloud-security-blog/11-advantages-of-cloud-computing-and-how-your-business-can-benefit-from-them/)
- [9] Internal sources (2018)
- [10] Techopedia Definition of Virtual Device (2018), Retrieved from [www.techopedia.com/definition/12632/virtual-device](http://www.techopedia.com/definition/12632/virtual-device)

- [11] OV (2018) - nVidia: nVidia Virtual GPU, Retrieved from [www.nvidia.com/en-us/design-visualization/solutions/virtualization/](http://www.nvidia.com/en-us/design-visualization/solutions/virtualization/)
- [12] OV (2018) - Kubernetes: What is Kubernetes, Retrieved from [kubernetes.io/docs/concepts/overview/what-is-kubernetes/](http://kubernetes.io/docs/concepts/overview/what-is-kubernetes/)
- [13] Dr. Thomas Fricke (01-16-2018): Kubernetes: Architektur und Einsatz – Eine Einführung mit Beispielen, Retrieved from [www.informatik-aktuell.de/entwicklung/methoden/kubernetes-architektur-und-einsatz-einfuehrung-mit-beispielen.html](http://www.informatik-aktuell.de/entwicklung/methoden/kubernetes-architektur-und-einsatz-einfuehrung-mit-beispielen.html)
- [14] OV (2018) - Docker: What is a container?, Retrieved from [www.docker.com/what-container](http://www.docker.com/what-container)
- [15] OV (2018) - CoreOS: Overview of a Pod, Retrieved from [coreos.com/kubernetes/docs/latest/pods.html](http://coreos.com/kubernetes/docs/latest/pods.html)
- [16] Jorge Acetozi (12-11-2017): Kubernetes Master Components: Etcd, API Server, Controller Manager, and Scheduler, Retrieved from [medium.com/jorgeacetozi/kubernetes-master-components-etcd-api-server-controller-manager-and-scheduler-3a0179fc8186](http://medium.com/jorgeacetozi/kubernetes-master-components-etcd-api-server-controller-manager-and-scheduler-3a0179fc8186)
- [17] Picture retrieved from [www.informatik-aktuell.de/fileadmin/\\_processed\\_/b/6/csm\\_K8s\\_abb1\\_fricke\\_b5ce6bd912.png](http://www.informatik-aktuell.de/fileadmin/_processed_/b/6/csm_K8s_abb1_fricke_b5ce6bd912.png)
- [18] OV (last viewed 08-05-2018) - deis: Kubernetes Overview, Retrieved from [deis.com/blog/2016/kubernetes-overview-pt-1/](http://deis.com/blog/2016/kubernetes-overview-pt-1/)
- [19] Christ Wright (22-06-2017): Kubernetes vs Docker Swarm, Retrieved from [platform9.com/blog/kubernetes-docker-swarm-compared](http://platform9.com/blog/kubernetes-docker-swarm-compared)
- [20] OV (last viewed 08-13-2018) - redhat: What is Kubernetes, Retrieved from [www.redhat.com/de/topics/containers/what-is-kubernetes](http://www.redhat.com/de/topics/containers/what-is-kubernetes)
- [21] OV (last viewed 08-13-2018) - Github.io: Running Spark on Kubernetes, Retrieved from [apache-spark-on-k8s.github.io/userdocs/running-on-kubernetes.html](http://apache-spark-on-k8s.github.io/userdocs/running-on-kubernetes.html)
- [22] OV (last viewed 08-13-2018): Apache Spark, Retrieved from [spark.apache.org/](http://spark.apache.org/)
- [23] Anzy Meerasahib (06-22-2018): Apache Spark and Big Data: What's Ahead, Retrieved from [tdwi.org/articles/2018/06/22/ta-all-apache-spark-and-big-data-whats-ahead.aspx](http://tdwi.org/articles/2018/06/22/ta-all-apache-spark-and-big-data-whats-ahead.aspx)

- [24] Image retrieved from [spark.apache.org/docs/latest/img/cluster-overview.png](https://spark.apache.org/docs/latest/img/cluster-overview.png)
- [25] OV (last viewed 08-08-2018) - Apache: Cluster Mode Overview, Retrieved from [spark.apache.org/docs/latest/cluster-overview.html](https://spark.apache.org/docs/latest/cluster-overview.html)
- [26] Ian Pointer (11-13-2017): What is Apache Spark? The big data analytics platform explained, Retrieved from [www.infoworld.com/article/3236869/analytics/what-is-apache-spark-the-big-data-analytics-platform-explained.html](http://www.infoworld.com/article/3236869/analytics/what-is-apache-spark-the-big-data-analytics-platform-explained.html)
- [27] OV (2018): Understanding Batch, Microbatch, and Streaming, Retrieved from [streaml.io/resources/tutorials/concepts/understanding-batch-microbatch-streaming](https://streaml.io/resources/tutorials/concepts/understanding-batch-microbatch-streaming)
- [28] Salman Salloum, Ruslan Dautov, Xiaojun Chen<sup>1</sup>, Patrick Xiaogang Peng, Joshua Zhexue Huang (07-20-2016): Big Data Analytics on Spark, published in Int J Data Sci Anal, Springer, Page 145-147
- [29] Ronald Barber, Christian Garcia-Arellan, Ronen Grosman, Rene Mueller, Vijayshankar Raman, Richard Sidle, Matt Spilchen, Adam Storm, Yuanyuan Tian, Pinar Tözün, Daniel Zilio, Matt Huras, Guy Lohman, C. Mohan, Fatma Özcan, Hamid Pirahesh (08-11-2017): Evolving Databases for New-Gen Big Data Applications, Retrieved from [researcher.watson.ibm.com/researcher/files/us-ytian/wildfire-cidr17.pdf](https://researcher.watson.ibm.com/researcher/files/us-ytian/wildfire-cidr17.pdf)
- [30] Chen Luo, Pinar Tözün, Yuanyuan Tian, Ronald Barber, Vijayshankar Raman, Richard Sidle (not published yet): Umzi: Unified MultiZone Indexing for LargeScale HTAP, Chapter 2, Internal Source
- [31] Madhura Maskasky (04-17-2018): Kubernetes Networking: Achieving High Performance with Calico, Retrieved from [platform9.com/blog/kubernetes-networking-achieving-high-performance-with-calico/](https://platform9.com/blog/kubernetes-networking-achieving-high-performance-with-calico/)

# Appendices